

The graph displays the KL divergence for ten agents, labeled 'Approx KL Agent 0' through 'Approx KL Agent 9'. The x-axis represents the 'Iteration' from 0 to 1000, and the y-axis represents the 'KL Divergence' from 0.00 to 0.04. All agents start with a KL divergence between 0.025 and 0.035 at iteration 0. They all show a general downward trend, with some initial fluctuations. By iteration 600, most agents have converged to a KL divergence between 0.01 and 0.015. After iteration 600, the values remain relatively stable, with a few minor peaks and troughs, indicating that the agents have reached a near-stationary state.

Iteration	Approx KL Agent 0	Approx KL Agent 1	Approx KL Agent 2	Approx KL Agent 3	Approx KL Agent 4	Approx KL Agent 5	Approx KL Agent 6	Approx KL Agent 7	Approx KL Agent 8	Approx KL Agent 9
0	0.030	0.028	0.025	0.032	0.027	0.026	0.029	0.028	0.027	0.028
200	0.018	0.017	0.016	0.019	0.017	0.016	0.018	0.017	0.016	0.017
400	0.012	0.011	0.010	0.013	0.011	0.010	0.012	0.011	0.010	0.011
600	0.011	0.010	0.009	0.012	0.010	0.009	0.011	0.010	0.009	0.010
800	0.010	0.009	0.008	0.011	0.009	0.008	0.010	0.009	0.008	0.009
1000	0.010	0.009	0.008	0.011	0.009	0.008	0.010	0.009	0.008	0.009

