

The graph displays the KL divergence for ten agents, labeled 'Approx KL Agent 0' through 'Approx KL Agent 9'. The x-axis represents iterations from 0.0 to 1.6, and the y-axis represents KL divergence on a logarithmic scale from 10^{-1} to 10^1 . All agents start with high KL divergence (between 10^0 and 10^1) and generally decrease over time. By iteration 0.8, most agents have converged to a KL divergence between 10^0 and 10^1 , with some agents showing more significant fluctuations than others.

