

The graph displays the training progress of 10 agents. The x-axis represents the number of episodes (0 to 1000), and the y-axis represents the total loss (0.0 to 1.0). The legend identifies the lines for each agent:

- Total Loss Agent 0 (Blue)
- Total Loss Agent 1 (Orange)
- Total Loss Agent 2 (Green)
- Total Loss Agent 3 (Red)
- Total Loss Agent 4 (Purple)
- Total Loss Agent 5 (Brown)
- Total Loss Agent 6 (Pink)
- Total Loss Agent 7 (Grey)
- Total Loss Agent 8 (Olive)
- Total Loss Agent 9 (Cyan)

Key observations from the graph include:

- Initial Loss:** At episode 0, Agent 0 has the highest loss (~0.8), while Agent 9 has the lowest (~0.4).
- Convergence:** By episode 1000, all agents' losses have converged to a similar range between 0.5 and 0.6.
- Stability:** Most agents show a relatively smooth decrease in loss, with some initial fluctuations. Agent 5 (Brown) shows a notable spike around episode 400.
- Final Performance:** Agent 9 (Cyan) maintains the lowest loss throughout most of the training, while Agent 0 (Blue) maintains the highest.

